

Topic 5: de novo assembly

Outcomes

- Identify the difference between de novo assembly and reference guided alignment
- Evaluate two different approaches to de novo genome assembly
- Describe how repetitive elements can hamper proper assembly and compare approaches that can overcome this problem
- Describe approaches for transcriptome/GBS de novo assembly


Introduction



stack of NY Times, June 27, 2000



stack of NY Times, June 27, 2000
on a pile of dynamite



this is just hypothetical



so, what did the June 27, 2000 NY
Times say?


Introduction

atshirt, approximately 6'2" 180 lbs. We have not yet named any suspects. Information is welcomed.

shirt, approximately 6'2" 180 lbs. We have not yet named any suspects. Information is welcomed. Please call

Introduction

Multiple identical
copies of a genome



Shatter the genome
into reads



Sequence the reads

AGAATATCA

TGAGAATAT

GAGAATATC

Assemble the
genome using
overlapping reads

AGAATATCA

GAGAATATC

TGAGAATAT

...TGAGAATATCA...

Alignment vs assembly


Aligning to a reference:

- Reference guided alignments: align the reads to a reference genome and looks for differences

Building a reference:

- *De novo* assembly: no previous genome assembly is used
- Comparative genome assembly: assemble a newly sequenced genome by mapping it on to a reference
- Hybrid approach: reference-guided and *de novo* for unused reads or *de novo* and then reference guided alignments

Introduction



Introduction

Original sequence

GATAGAAGGGTCCGCTCGCTCAGCTACCGGTTTTTATAGATCTA

GATAGAAGGGTCCGCT
AGAAGGGTCCGCTC
GGGTCCGCTCGCTCA
CCGCTCGCTCAGC
CTCGCTCAGCTACC
TCAGCTACCGGTTT
CTACCGGTTTTT
AGCTACCGGTTTTTAT
TTTTTATAGATCTA



fragmented sequences
from sequencer
(reads)

Introduction

assembled

fragmented sequences
from sequencer
(reads)

TTTTTATAGATCTA

AGCTACCGGTTTTTAT

CAGCTACCGGTTTTT

TCAGCTACCGGTTT

CTCGCTCAGCTACC

CCGCTCGCTCAGC

GGGTCCGCTCGCTCA

AGAAGGGTCCGCTC

GATAGAAGGGTCCGCT

GATAGAAGGGTCCGCTCGCTCAGCTACCGGTTTTTATAGATCTA

We want to reconstruct this from the reads

Introduction

Simplified scenario

- Single strand
- Error free
- Complete coverage

TTTTTATAGATCTA

AGCTACCGGTTTTTAT

CAGCTACCGGTTTTT

TCAGCTACCGGTTT

CTCGCTCAGCTACC

CCGCTCGCTCAGC

GGGTCCGCTCGCTCA

AGAAGGGTCCGCTC

GATAGAAGGGTCCGCT

GATAGAAGGGTCCGCTCGCTCAGCTACCGGTTTTTATAGATCTA

Introduction

Coverage: reads “covering” a position in the genome
(average or at a single base or region)

TTTTTATAGATCTA
AGCTACCGGTTTTTAT
CAGCTACCGGTTTTT
TCAGCTACCGGTTT
CTCGCTCAGCTACC
CCGCTCGCTCAGC
GGGTCCGATCGCTCA
AGAAGGGTCCGCTC
GATAGAAGGGTCCGCT

131 bases in the reads

↓

44 bases in the “genome”

GATAGAAGGGTCCGCTCGCTCAGCTACCGGTTTTTATAGATCTA

What is our average coverage?

What is the coverage at the arrow?

Introduction

TTTTTATAGATCTA
AGCTACCGGTTTTTAT
CAGCTACCGGTTTTT
TCAGGTACCGGTTT
CTCGCTCAGCTACC
CCGCTCGCTCAGC
GGGTCCGATTGCTCA
AGAAGGGTCCGCTC
GATAGAAGGGTCCGCT

GATAGAAGGGTCCGCTCGCTCAGCTACCGGTTTTTATAGATCTA

Why might there be differences among reads
covering the same position?

Introduction

CCGCTCGCTCAGC
TCAGCTACCGGTTT
CTCGCTCAGCTACC
CAGCTACCGGTTTTT
AGAAGGGTCCGCTC
GATAGAAGGGTCCGCT
AGCTACCGGTTTTTAT
TTTTTATAGATCTA
GGGTCCGCTCGCTCA

How would you go about “assembling” these reads when you have no reference?

Code break

Write some code to find all the overlaps exactly 4 bp in length between **CTCTAGGCC** and a list of other sequences in the file `~/Topic_5/data/overlaps.fa`

Overlap-layout-consensus

Overlap: make an overlap graph

Layout: find the path through the graph

Consensus: find the most likely contig sequence


OLC programs:

ARACHNE, PHRAP, CAP, TIGR, CELERA

Overlap-layout-consensus

Reads: CTCTAGGCC GCCTCAAT CAATTTT

This is a
graph!(directed)



Node or vertex=
read

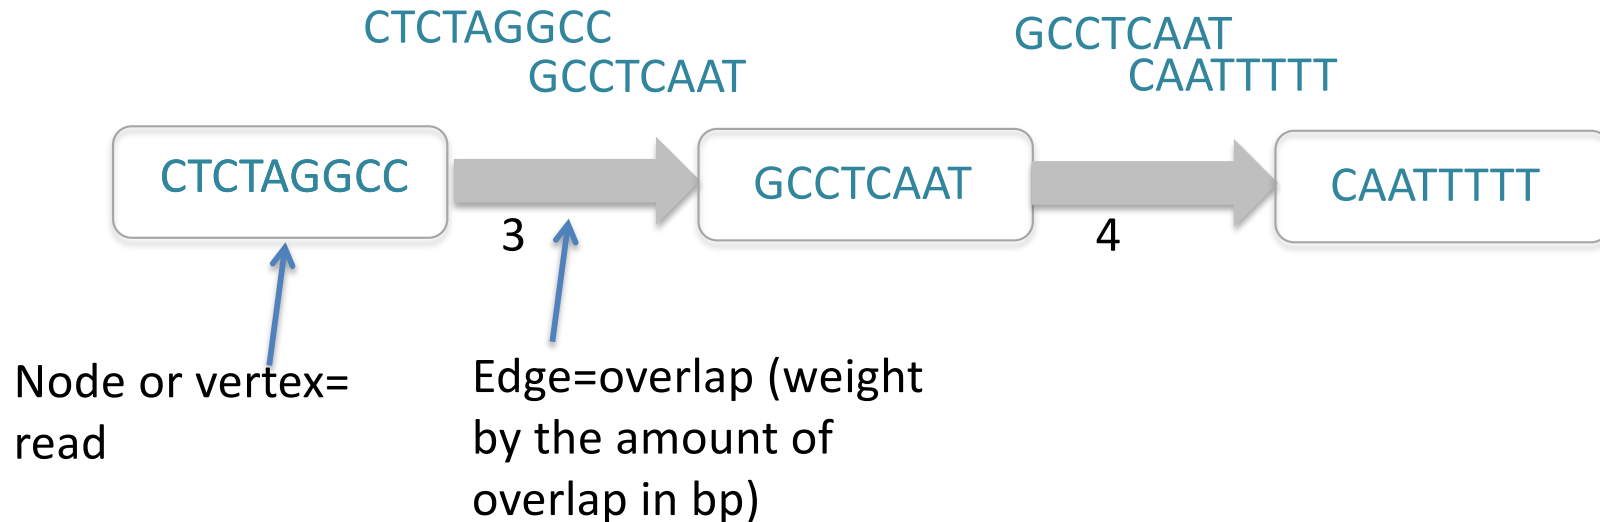
Edge=overlap (weight
by the amount of
overlap in bp)

Can pick a minimum overlap length (e.g. 3 bp)

Finding overlaps can be computationally challenging
when you have millions of reads!

Overlap-**layout**-consensus


Reads: CTCTAGGCC GCCTCAAT CAATTTT



Here we have only one path through the graph

Overlap-**layout**-consensus

These graphs get complicated!



Minimum
overlap = 3
Read length = 7


GCATTATATATTGCGCGTACGGCGCCGCTACA

Original sequence

How can we find the best path?

Overlap-**layout**-consensus

Hamiltonian path: hit each node (read) once
–no quick way to figure it out (NP-complete)
–not practical and not implemented



Overlap-**layout**-consensus

Shortest superstring: find the shortest final sequence (greatest overlap between reads)


- hit each node (read) once

- NP-hard

Overlap-**layout**-consensus

Greedy algorithm (example)

- 1) Pairwise alignments between all fragments
- 2) Pick the two with the largest overlap
- 3) Merge chosen fragments
- 4) Repeat



the greedy one

Overlap-layout-**consensus**

Join sequences together into one sequence

Reads: CTCTAGGCC GCCTCAAT CAATTTT

CTCTAGGCC
GCCTCAAT

GCCTCAAT
CAATTTT

CTCTAGGCC

3

GCCTCAAT

4

CAATTTT

CTCTAGGCC
GCCTCAAT
CAATTTT

CTCTAGGCCCTCAATTTT



Overlap-layout-**consensus**

Limitations of OLC

- require overlaps to be scored between all possible pairs of reads. This is a problem when you have millions of reads
- finding the best path through the graph with a huge number of nodes (reads) is computationally challenging

Is there a faster way to assemble many short reads?

De Bruijn graphs

What are all the 5-mers (5 bp fragments) in these reads?

2 reads of 9 bp

read 1

ATGGGGAAC

ATGGG

TGGGG

GGGGA

GGGAA

GGAAC

read 2

GGGAACCCC

GGGAA

GGAAC

GAACC

AACCC

ACCCC

If a read is L bp long, how many kmers of size k can you make?

Code break

Find all the unique 9mers in a fasta sequence and sort them alphabetically `~/Topic_5/data/kmer.fa`

1. Find all the kmers in this fasta sequence.

Hints: test out the following commands.

The cut program lets you pull out columns of text

```
cut -c2- kmer.fa
```

```
cut -c1-4 kmer.fa
```

```
for num in {1..10}
```

```
do
```

```
echo $num >> file.txt
```

```
done
```

2. Sort them and keep the unique ones

Hint: try `sort`

De Bruijn graphs

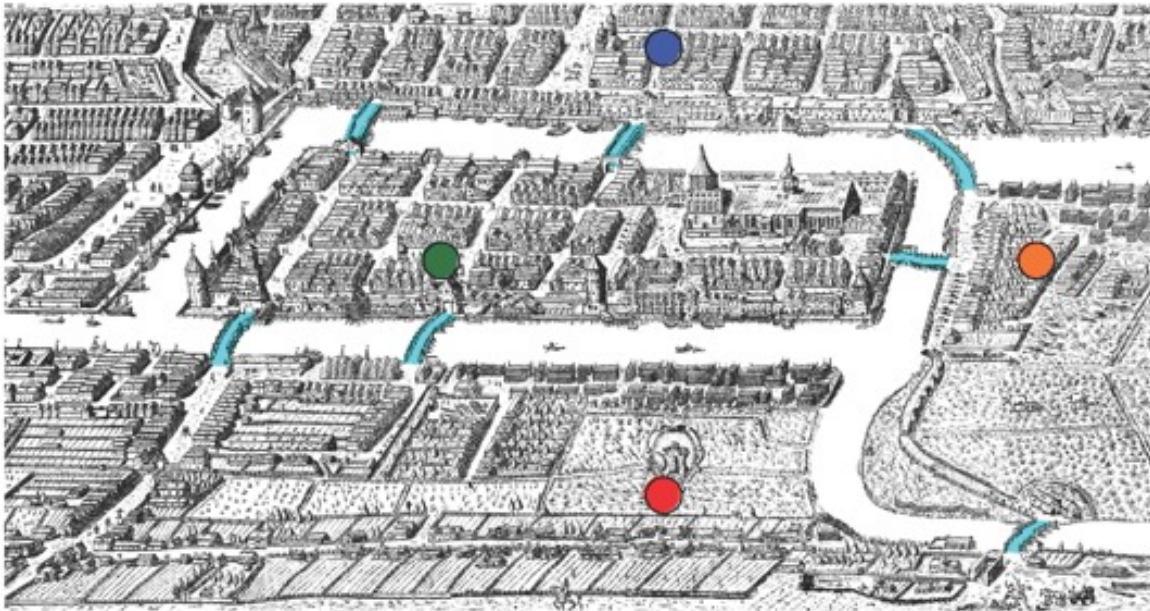
- Join up all the k-mers (length = k bp) into a graph with an overlap of k-1 (here k=5)




- Traverse through the graph
- The first base of each node spells out the sequence

De Bruijn graphs

a





b




De Bruijn graphs

Eulerian graph must be both balanced and strongly connected




De Bruijn graphs




Algorithm to find a path through an Eulerian graph

De Bruijn graphs





Algorithm to find a path through an Eulerian graph

De Bruijn graphs



Algorithm to find a path through an Eulerian graph

De Bruijn graphs



Algorithm to find a path through an Eulerian graph

De Bruijn graphs

Limitations of the Eulerian path:

- With “perfect” genomic data there are usually many Eulerian tours
- Data is not perfect (areas of low coverage, errors, repeats, etc.)

De Bruijn graphs

TAGTCGAGGCTTTAGATCCGATGAGGCTTTAGAGACAG


1. Sequence

AGTCGAG	CTTTAGA	CGATGAG	CTTTAGA
GTCGGG	TTAGATC	ATGAGGC	GAGACAG
GAGGCTC	ATCCGAT	AGGCTTT	GAGACAG
AGTCGAG	TAGATCC	ATGAGGC	TAGAGAA
TAGTCGA	CTTTAGA	CCGATGA	TTAGAGA
CGAGGCT	AGATCCG	TGAGGCT	AGAGACA
TAGTCGA	GCTTTAG	TCCGATG	GCTCTAG
TCGACGC	GATCCGA	GAGGCTT	AGAGACA
TAGTCGA	TTAGATC	GATGAGG	TTTAGAG
GTCGAGG	TCTAGAT	ATGAGGC	TAGAGAC
AGGCTTT	ATCCGAT	AGGCTTT	GAGACAG
AGTCGAG	TTAGATT	ATGAGGC	AGAGACA
GGCTTTA	TCCGATG	TTTAGAG	
CGAGGCT	TAGATCC	TGAGGCT	GAGACAG
AGTCGAG	TTTAGATC	ATGAGGC	TTAGAGA
GAGGCTT	GATCCGA	GAGGCTT	GAGACAG

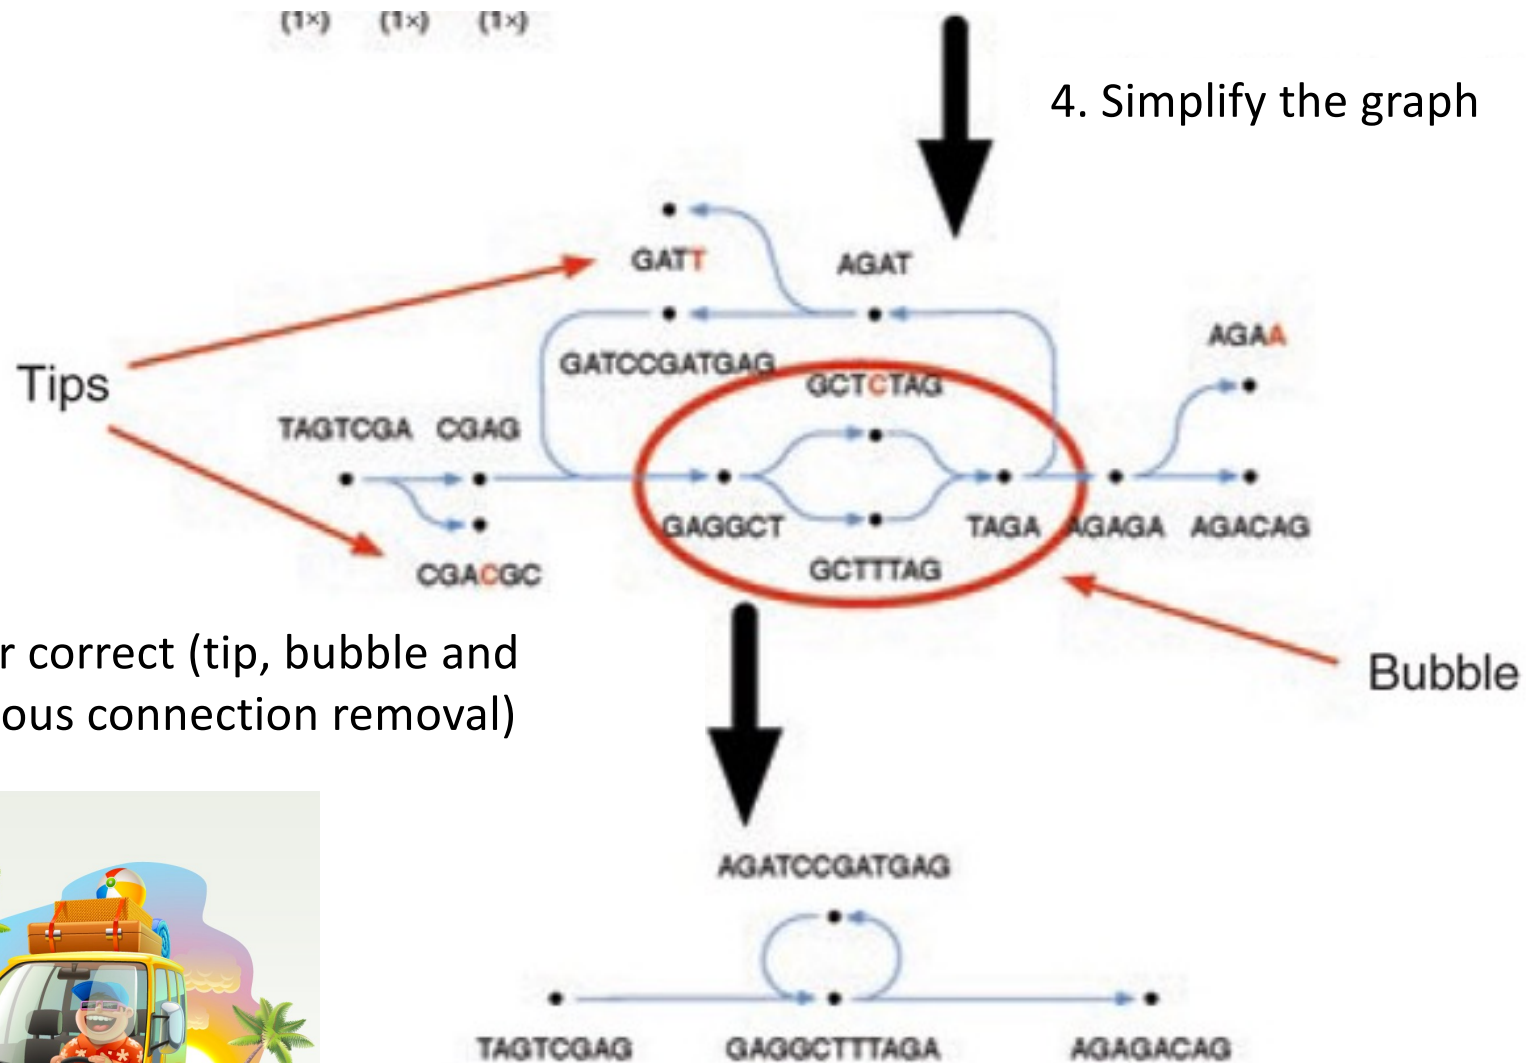
Sequencing errors

2. Find the kmers

3. Build the graph



De Bruijn graphs



De Bruijn graphs

Advantages:

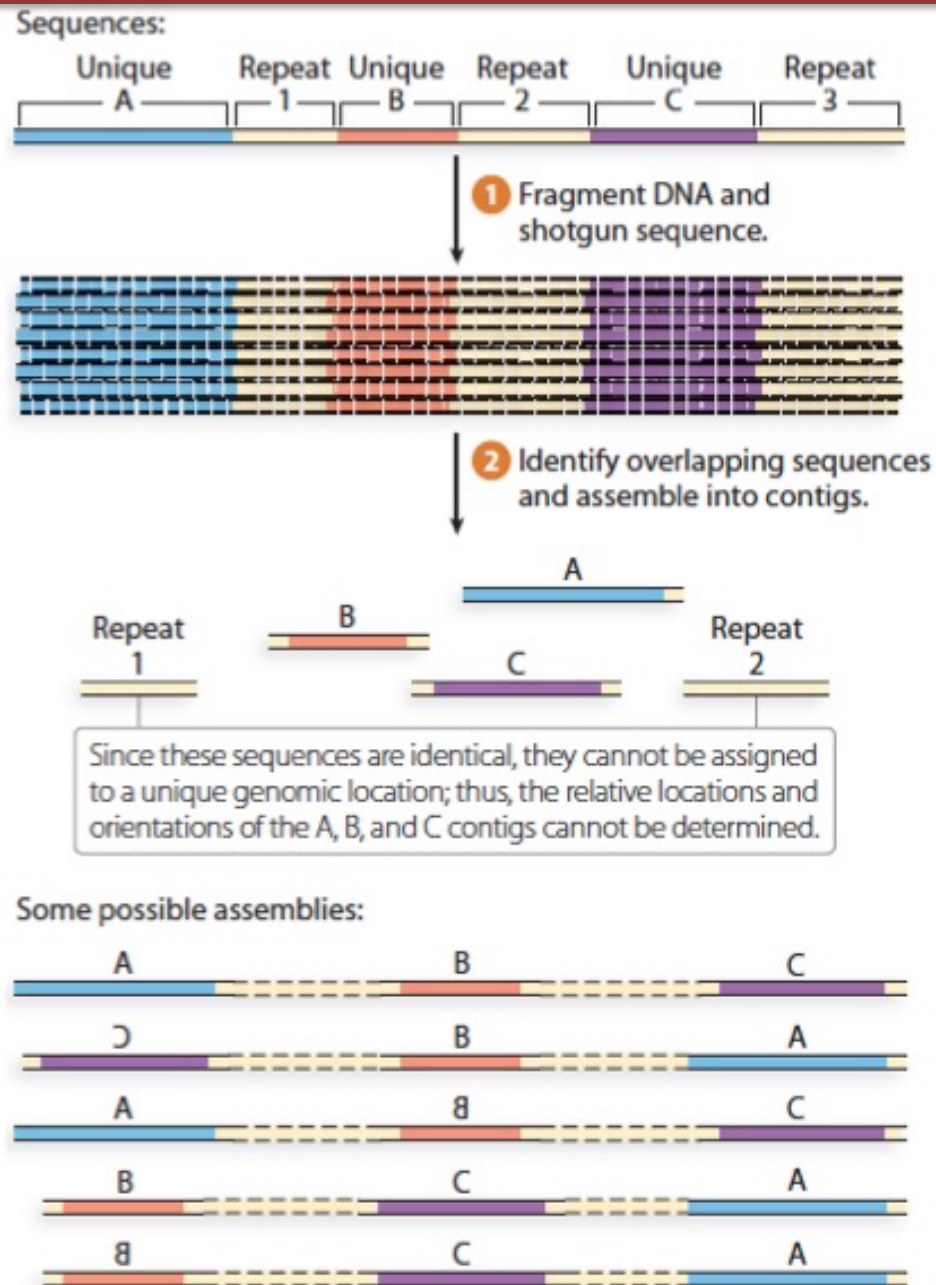
- 1) Set node length (no overlap algorithm)
- 2) Easy approaches for traversing through the graph
- 3) Simpler representation of repeats in the graph

Disadvantages:


- 1) Lose information
- 2) Shorter contigs

For PacBio and other long read sequences, what type of assembly strategy would you use?

Repetitive regions




Repetitive regions




Repetitive regions

- Finishing eukaryotic genome assemblies can be challenging because much of the genome is repetitive
- This repetitive DNA breaks up the assembly and obscures the order and orientation of the assembled contigs
- Even well studied model organisms can have poorly assembled regions of their genome

Repetitive regions




Repetitive regions



Current assembly approaches

- Long read sequencing
- Synthetic long reads
- Long-range scaffolding technologies



Long read assemblies

Long read only *de novo* assembly. PacBio/Nanopore reads are assembled using an OLC algorithm (e.g., HGAP).(>50x PacBio, or 15X PacBio HiFi)

Hybrid *de novo* assembly. Error correct long reads with more accurate short reads (e.g., PacBioToCA module of Celera) before performing long read assembly. (~20x PacBio)

Gap filling. Starting with an *existing* mate-pair based assembly, the internal gaps (consisting of Ns) inside the scaffolds are filled using PacBio sequences. (~5x PacBio)

Scaffolding. Using an *existing* assembly (such as an assembly based on short read data), PacBio reads are used to join contigs. (~5x PacBio)

Synthetic long read assemblies

Synthetic long reads (SLRs) technologies [Illumina, 10X Genomics, Loop Genomics, and Universal Sequencing Technology (UST)]




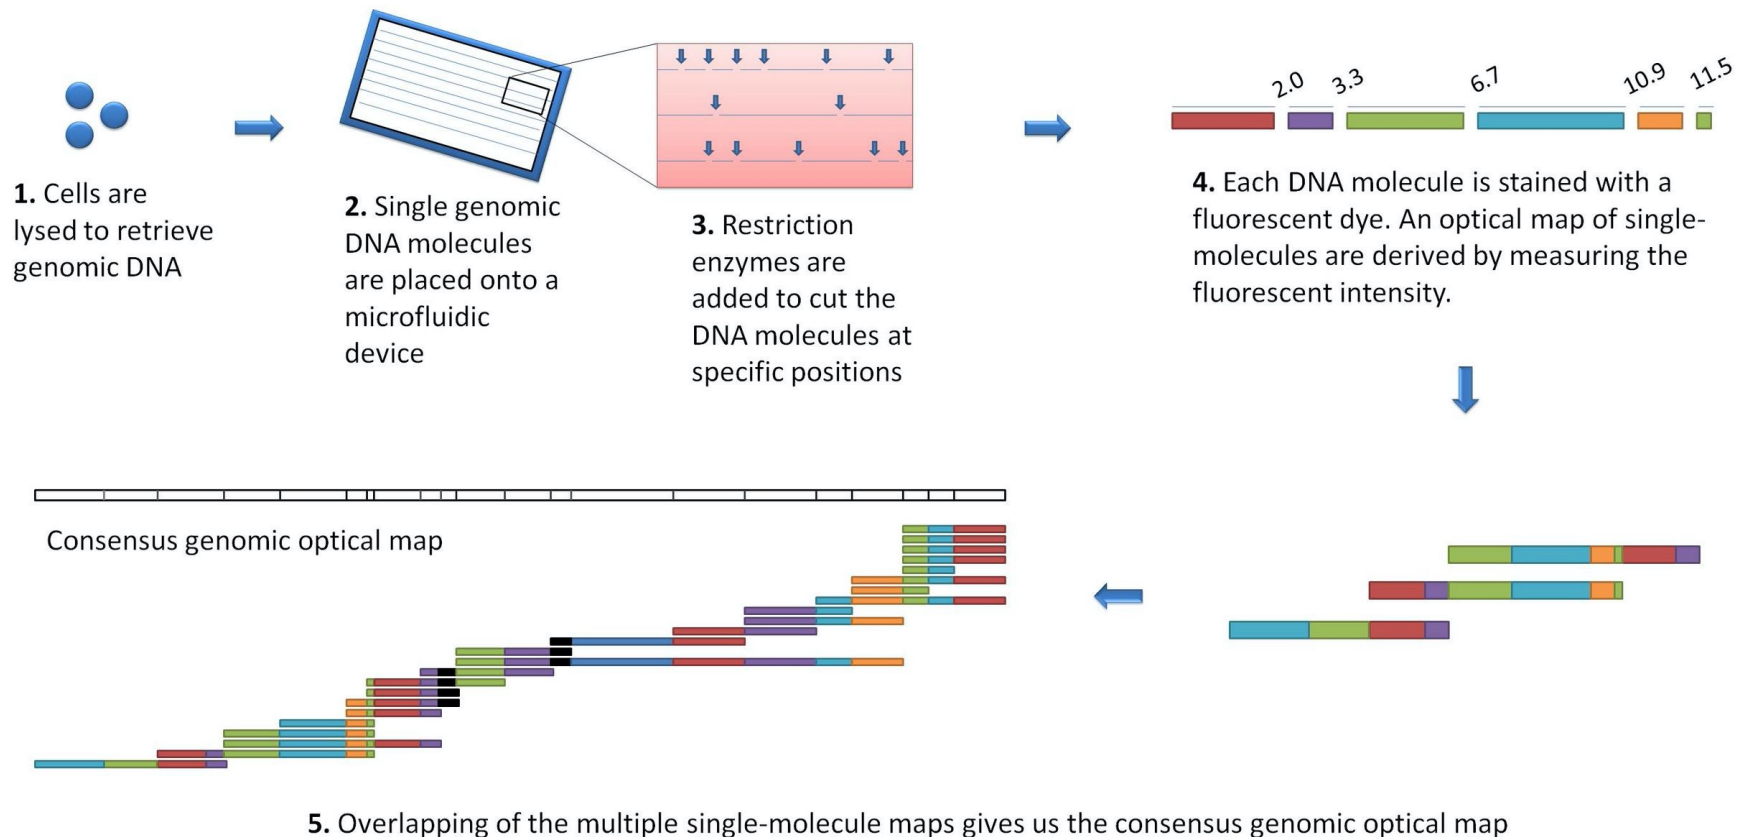
Figure 1 | The TSLR technology. The barcode assembly step generates virtual long reads. In an idealized scenario, the barcode assembly would result in ~300 TSLRs with lengths of ~10 kb. In reality, it results in 350–450 TSLRs varying in length from 1 to 10 kb.

Anton Bankevich, & Pavel A Pevzner. (2016). TruSPAdes: Barcode assembly of TruSeq synthetic long reads. *Nature Methods*, 13(3), 248-250.

Long-range scaffolding technologies

Optical mapping

- Bionano Genomics <https://vimeo.com/116090215>




Long-range scaffolding technologies

Omni-C (similar to Hi-C) is one approach to get chromosome level assemblies

Paired reads with small and big distances between them are created. There are more paired reads close together than far apart.

These reads are aligned to the contigs
Contigs with more joins (paired reads aligned across two contigs and therefore “joined”) are closer together in the genome than those with few joins

<https://www.youtube.com/watch?v=-MxEw3IXUWU>



Long-range scaffolding technologies

HiRise Scaffolding Report

December 15, 2020

Cakile edentula


Kay Hodgins

Monash University



Contents

- [Overview](#)
- [Contiguity Metrics](#)
- [BUSCO](#)
- [Pair Size Distribution](#)
- [Scaffolding Summary](#)



Overview

Assembly	Total Length (bp)	N50	L50	N90	L90
Input Assembly	651,503,399	1,397,589	115	309,490	501
Dovetail HiRise Assembly	651,583,577	68,669,067	5	54,597,944	9

Other types of de novo assembly

Transcriptome

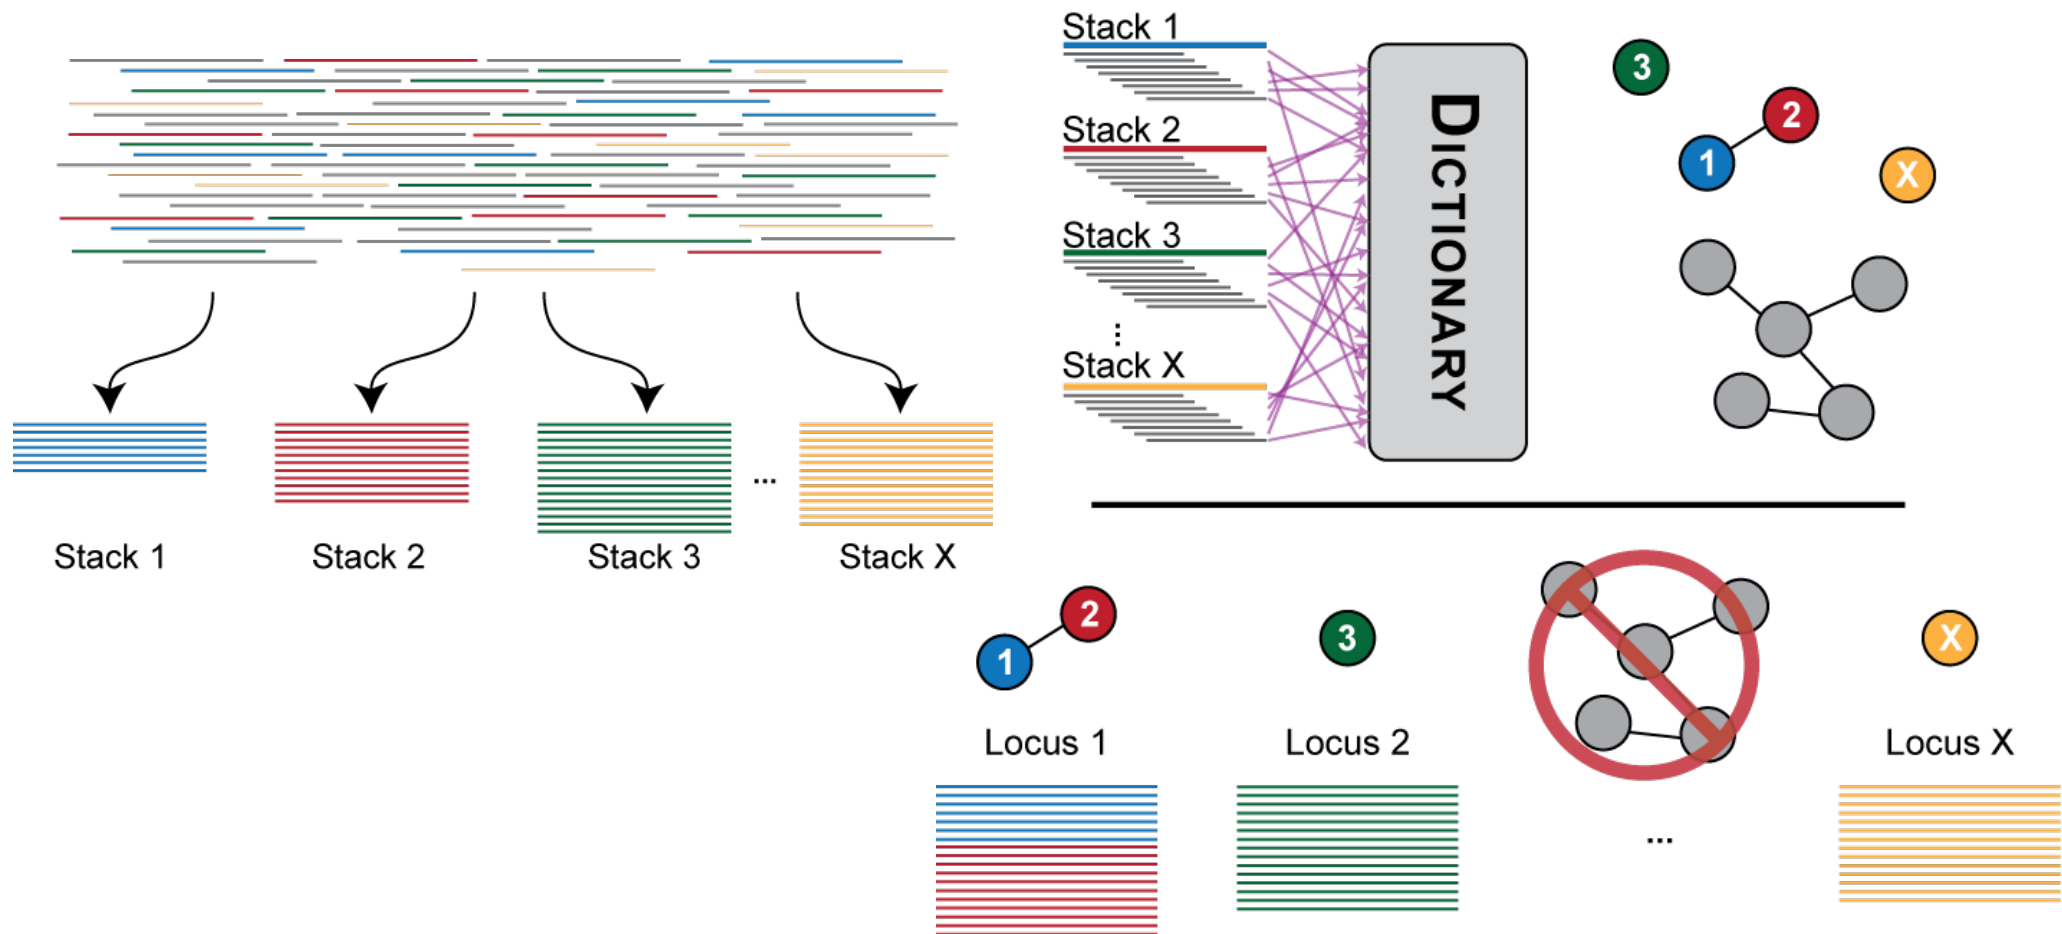
- Variable coverage among genes/isoforms
- Alternative splicing
promoters, exons, and poly(A)

PacBio/Nanopore long reads – full length
isoforms

Illumina short reads – Trinity (recommended)

Other types of de novo assembly

De novo assembly of GBS reads: Stacks



Further Reading

Jiao WB, Schneeberger K. 2017. The impact of third generation genomic technologies on plant genome assembly. *Current Opinion in Plant Biology* 36: 64–70.

Flicek, P., & Birney, E. (2009). Sense from sequence reads: methods for alignment and assembly. *Nature methods*, 6, S6-S12.

Zerbino DR, Birney E. Velvet: Algorithms for de novo short read assembly using de Bruijn graphs. *Genome Research*. 2008;18(5):821-829. doi:10.1101/gr.074492.107.

<http://computing.bio.cam.ac.uk/local/doc/velvet.pdf>

Li, Z., Chen, Y., Mu, D., Yuan, J., Shi, Y., Zhang, H., ... & Yang, B. (2012). Comparison of the two major classes of assembly algorithms: overlap–layout–consensus and de-bruijn-graph. *Briefings in functional genomics*, 11(1), 25-37.

Grabherr MG, Haas BJ, Yassour M, et al. Trinity: reconstructing a full-length transcriptome without a genome from RNA-Seq data. *Nature biotechnology*. 2011;29(7):644-652. doi:10.1038/nbt.1883.

<https://github.com/trinityrnaseq/trinityrnaseq/wiki>

J. Catchen, A. Amores, P. Hohenlohe, W. Cresko, and J. Postlethwait. Stacks: building and genotyping loci de novo from short-read sequences. *G3: Genes, Genomes, Genetics*, 1:171-182, 2011.

<http://catchenlab.life.illinois.edu/stacks/>

Tutorial: short read assembly

Today you will assemble a bacterial genome using short reads and long reads and then compare the assemblies

Velvet overview (short reads):

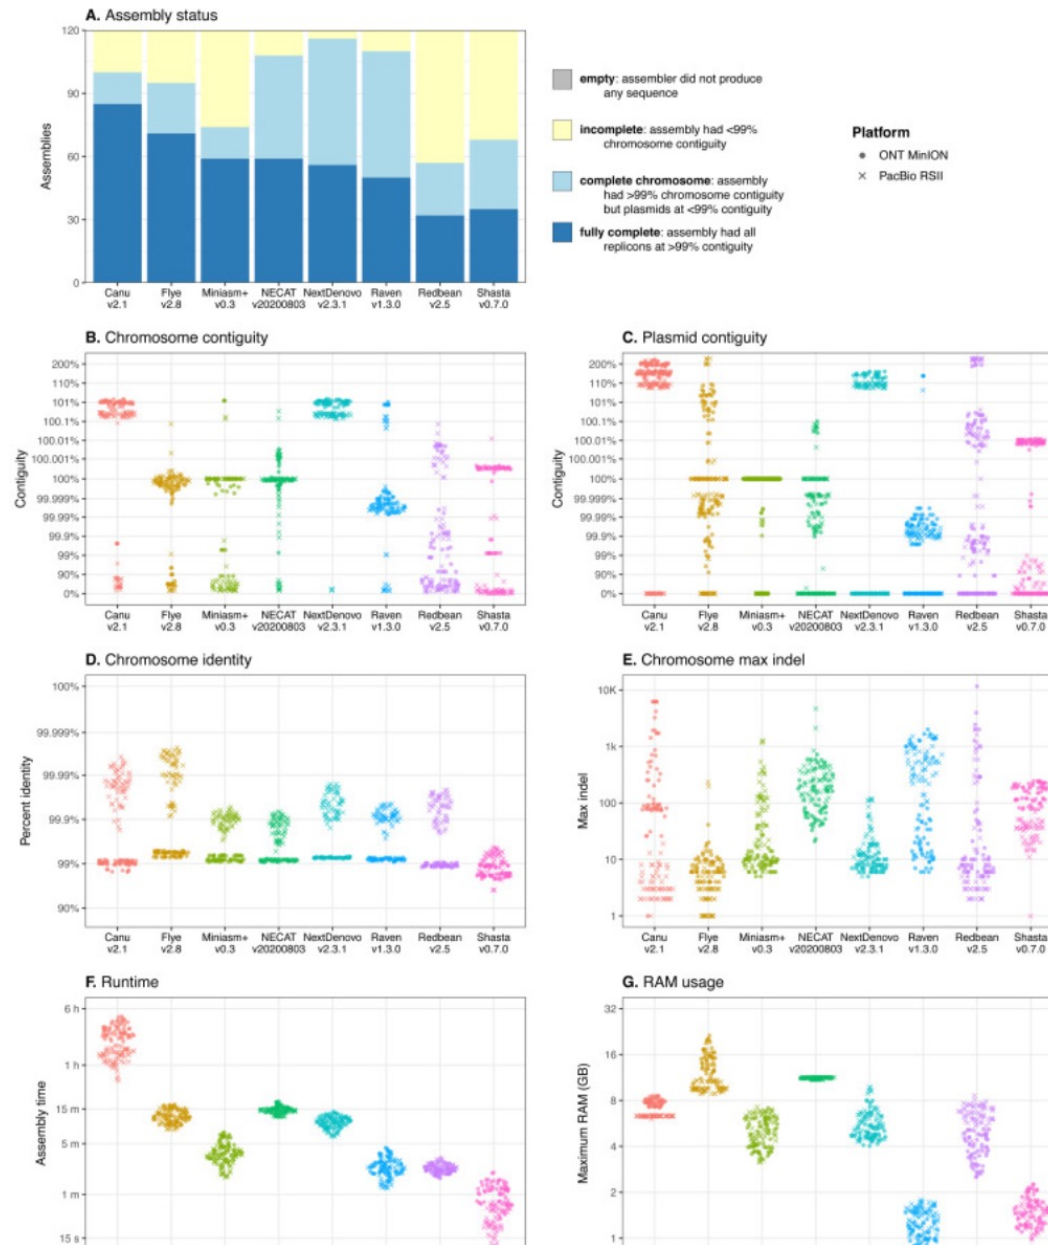
1. Hash k-mers
2. Construct the graph
3. Correct for errors
4. Resolve the repeats

Minimap2/Miniasm/Racon (long reads):

1. Minimap2 to map the reads to each other (overlap)
2. Miniasm to assemble (layout)
3. Minimap2/Racon (consensus)

Long read assemblers benchmarking

Real read set assembly results



Tutorial: assembly metrics

Finally, we need to assess the quality of the assemblies

There are a number of metrics. Below are a few examples:

Final assembly length (is it close to the expected size?)

N50 (50% of the genome is in a contig of that size or larger)

L50 (as the smallest number of contigs whose length sum makes up half of genome size)

The percentage of Ns

We are going to use the program Quast

We are also going to view the assembly graphs using Bandage

Tutorial: assembly metrics

The number of conserved single-copy orthologs (e.g., BUSCO) that are complete, fragmented or duplicated

